

# clEnqueueNDRangeKernel

Enqueues a command to execute a kernel on a device.

```
cl_int clEnqueueNDRangeKernel (cl_command_queue command_queue,
                                cl_kernel kernel,
                                cl_uint work_dim,
                                const size_t *global_work_offset,
                                const size_t *global_work_size,
                                const size_t *local_work_size,
                                cl_uint num_events_in_wait_list,
                                const cl_event *event_wait_list,
                                cl_event *event)
```

## Parameters

*command\_queue*

A valid command-queue. The kernel will be queued for execution on the device associated with *command\_queue*.

*kernel*

A valid kernel object. The OpenCL context associated with *kernel* and *command\_queue* must be the same.

*work\_dim*

The number of dimensions used to specify the global work-items and work-items in the work-group. *work\_dim* must be greater than zero and less than or equal to three.

*global\_work\_offset*

Must currently be a NULL value. In a future revision of OpenCL, *global\_work\_offset* can be used to specify an array of *work\_dim* unsigned values that describe the offset used to calculate the global ID of a work-item instead of having the global IDs always start at offset (0, 0,... 0).

*global\_work\_size*

Points to an array of *work\_dim* unsigned values that describe the number of global work-items in *work\_dim* dimensions that will execute the kernel function. The total number of global work-items is computed as *global\_work\_size*[0] \*...\* *global\_work\_size*[*work\_dim* - 1].

The values specified in *global\_work\_size* cannot exceed the range given by the `sizeof(size_t)` for the device on which the kernel execution will be enqueued. The `sizeof(size_t)` for a device can be determined using `CL_DEVICE_ADDRESS_BITS` in the table of OpenCL Device Queries for

[clGetDeviceInfo](#). If, for example, `CL_DEVICE_ADDRESS_BITS = 32`, i.e. the device uses a 32-bit address space, `size_t` is a 32-bit unsigned integer and `global_work_size` values must be in the range  $1 \dots 2^{32} - 1$ . Values outside this range return a `CL_OUT_OF_RESOURCES` error.

### *local\_work\_size*

Points to an array of *work\_dim* unsigned values that describe the number of work-items that make up a work-group (also referred to as the size of the work-group) that will execute the kernel specified by *kernel*. The total number of work-items in a work-group is computed as `local_work_size[0] * ... * local_work_size[work_dim - 1]`. The total number of work-items in the work-group must be less than or equal to the

`CL_DEVICE_MAX_WORK_GROUP_SIZE` value specified in table of OpenCL Device Queries for [clGetDeviceInfo](#) and the number of work-items specified in `local_work_size[0], ... local_work_size[work_dim - 1]` must be less than or equal to the corresponding values specified by

`CL_DEVICE_MAX_WORK_ITEM_SIZES[0], ...`

`CL_DEVICE_MAX_WORK_ITEM_SIZES[work_dim - 1]`. The explicitly specified *local\_work\_size* will be used to determine how to break the global work-items specified by *global\_work\_size* into appropriate work-group instances. If *local\_work\_size* is specified, the values specified in `global_work_size[0], ... global_work_size[work_dim - 1]` must be evenly divisible by the corresponding values specified in `local_work_size[0], ... local_work_size[work_dim - 1]`.

The work-group size to be used for *kernel* can also be specified in the program source using the `__attribute__((reqd_work_group_size(X, Y, Z)))` qualifier. In this case the size of work group specified by *local\_work\_size* must match the value specified by the `reqd_work_group_size` `__attribute__` qualifier.

*local\_work\_size* can also be a NULL value in which case the OpenCL implementation will determine how to break the global work-items into appropriate work-group instances.

See note for more information.

### *event\_wait\_list* and *num\_events\_in\_wait\_list*

Specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same.

### *event*

Returns an event object that identifies this particular kernel execution instance. Event objects are unique and can be used to identify a particular

kernel execution instance later on. If *event* is NULL, no event will be created for this kernel execution instance and therefore it will not be possible for the application to query or queue a wait for this particular kernel execution instance.

## Notes

Work-group instances are executed in parallel across multiple compute units or concurrently on the same compute unit.

Each work-item is uniquely identified by a global identifier. The global ID, which can be read inside the kernel, is computed using the value given by *global\_work\_size* and *global\_work\_offset*. In OpenCL 1.0, the starting global ID is always (0, 0, ... 0). In addition, a work-item is also identified within a work-group by a unique local ID. The local ID, which can also be read by the kernel, is computed using the value given by *local\_work\_size*. The starting local ID is always (0, 0, ... 0).

## Errors

Returns CL\_SUCCESS if the kernel execution was successfully queued. Otherwise, it returns one of the following errors:

- CL\_INVALID\_PROGRAM\_EXECUTABLE if there is no successfully built program executable available for device associated with *command\_queue*.
- CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid command-queue.
- CL\_INVALID\_KERNEL if *kernel* is not a valid kernel object.
- CL\_INVALID\_CONTEXT if context associated with *command\_queue* and *kernel* is not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- CL\_INVALID\_KERNEL\_ARGS if the kernel argument values have not been specified.
- CL\_INVALID\_WORK\_DIMENSION if *work\_dim* is not a valid value (i.e. a value between 1 and 3).
- CL\_INVALID\_WORK\_GROUP\_SIZE if *local\_work\_size* is specified and number of work-items specified by *global\_work\_size* is not evenly divisible by size of work-group given by *local\_work\_size* or does not match the work-group size specified for *kernel* using the [\\_\\_attribute\\_\\_\(\(reqd\\_work\\_group\\_size\(X, Y, Z\)\)\)](#) qualifier in program source.
- CL\_INVALID\_WORK\_GROUP\_SIZE if *local\_work\_size* is specified and the total number of work-items in the work-group computed as *local\_work\_size*[0] \* ... \* *local\_work\_size*[*work\_dim* - 1] is greater than the

value specified by `CL_DEVICE_MAX_WORK_GROUP_SIZE` in the table of OpenCL Device Queries for [clGetDeviceInfo](#).

- `CL_INVALID_WORK_GROUP_SIZE` if *local\_work\_size* is NULL and the [\\_\\_attribute\\_\\_\(\(reqd\\_work\\_group\\_size\(X, Y, Z\)\)\)](#) qualifier is used to declare the work-group size for *kernel* in the program source.
- `CL_INVALID_WORK_ITEM_SIZE` if the number of work-items specified in any of *local\_work\_size*[0], ... *local\_work\_size*[*work\_dim* - 1] is greater than the corresponding values specified by `CL_DEVICE_MAX_WORK_ITEM_SIZES`[0], ..., `CL_DEVICE_MAX_WORK_ITEM_SIZES`[*work\_dim* - 1].
- `CL_INVALID_GLOBAL_OFFSET` if *global\_work\_offset* is not NULL.
- `CL_OUT_OF_RESOURCES` if there is a failure to queue the execution instance of *kernel* on the command-queue because of insufficient resources needed to execute the kernel. For example, the explicitly specified *local\_work\_size* causes a failure to execute the kernel because of insufficient resources such as registers or local memory. Another example would be the number of read-only image args used in *kernel* exceed the `CL_DEVICE_MAX_READ_IMAGE_ARGS` value for device or the number of write-only image args used in *kernel* exceed the `CL_DEVICE_MAX_WRITE_IMAGE_ARGS` value for device or the number of samplers used in *kernel* exceed `CL_DEVICE_MAX_SAMPLERS` for device.
- `CL_MEM_OBJECT_ALLOCATION_FAILURE` if there is a failure to allocate memory for data store associated with image or buffer objects specified as arguments to *kernel*.
- `CL_INVALID_EVENT_WAIT_LIST` if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- `CL_OUT_OF_HOST_MEMORY` if there is a failure to allocate resources required by the OpenCL implementation on the host.

## Specification

 [OpenCL Specification](#)

## Also see

[clCreateCommandQueue](#), [clGetDeviceInfo](#), [clEnqueueNativeKernel](#), [clEnqueueTask](#), [Work-Item Functions](#)

Copyright © 2007-2009 The Khronos Group Inc. Permission is hereby granted, free of charge, to any person obtaining a copy of this software and/or associated documentation files (the "Materials"), to deal in the Materials without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Materials, and to permit persons to whom the Materials are furnished to do so, subject to the condition that this copyright notice and permission notice shall be included in all copies or substantial portions of the Materials.